

Website audit and remediation plan

cyberskill.world

Prepared for: Stephen Cheng (Trinh Thai Anh), Founder

CyberSkill Software Solutions Consultancy and Development JSC, Ho Chi Minh City

Assessment date: 11 July 2026

Scope: live audit and benchmark of the public website, with a phased plan for global client and partner acquisition

Executive summary

The CyberSkill site is a modern, well-engineered marketing site with real art direction and a technical foundation that already sits ahead of most boutique software consultancy sites. It is built on Next.js, server-renders its content, carries a distinctive dark and gold identity anchored by the Lumi genie mascot, and runs a coherent story from a clear wish to working software. Its structured data, bilingual setup, and on-page basics are strong.

The gaps that hold it back from converting attention into clients and partners are not about polish. They are about proof, measured performance, and a short list of precise search and answer-engine fixes. Case studies are anonymized with no client names or numbers, there is no real team page, and there are no logos, testimonials, or certifications, which are the exact signals international and enterprise buyers look for. The homepage carries roughly 0.74 MB of JavaScript and a constant WebGL animation with no verified Core Web Vitals, so responsiveness on lower-end phones is a risk that is currently unmeasured. A few small technical items, such as run-together animated headings, untranslated Vietnamese page titles, and a missing llms.txt, are quick to fix and worth doing.

Overall the site rates 78 out of 100, a strong B-plus. The design and build quality are a genuine differentiator. The work now is to put verifiable proof, measured performance, and a handful of SEO and GEO fixes behind that quality so the site earns and closes the business the founder wants to grow.

Headline strengths

- Rich, correct structured data: Organization with legal name, founding date, address, phone, and a DUNS identifier, plus FAQPage and a Service catalog.
- Clean technical SEO: unique title and meta description, self-referencing canonical, index-follow, a valid robots.txt and a 20-URL bilingual sitemap.
- Reciprocal hreflang for English and Vietnamese, with a genuinely translated Vietnamese site, which suits the global-growth goal.
- Distinctive, cohesive brand and confident visual design that stands apart from template-driven competitors.
- Fast edge delivery (server response near 65 to 75 ms) and accessibility groundwork such as a skip link, labelled contact form, and reduced-motion CSS.

Highest-impact gaps

- Credibility proof is missing: anonymized case studies with no metrics, no named team, no client logos, testimonials, third-party ratings, or certifications.
- Performance is unverified and at risk: heavy JavaScript and always-on animation, no field Core Web Vitals data, and Lighthouse could not be pulled during this audit.
- A few precise SEO and GEO fixes: animated homepage headings serialize without word spaces, Vietnamese page titles are not localized, and there is no llms.txt or deliberate AI-crawler stance.
- No authority content: there is no blog or insights section to rank, to be cited by AI answer engines, or to demonstrate expertise.

Scorecard

Each dimension is scored out of 100 with a letter grade and a one-line verdict. Scores reflect the live evidence gathered on 11 July 2026 and are judged against current best practice for software consultancy sites, described in the benchmark section.

Dimension	Score	Grade	Verdict
Technical SEO	88	A-	Excellent structured data, meta, hreflang, and sitemap; minor localization and rendering nuances.
On-page SEO	80	B+	Good titles, description, and heading order; animated headings lose word spaces, some thin crawlable copy.
GEO (AI answer engines)	72	B-	Strong schema, FAQ, and entity clarity; add llms.txt, quotable content, and an explicit AI-crawler stance.
Performance and Core Web Vitals	66	B-	Fast edge and server render, but heavy JavaScript and constant animation, with no verified field data.
UI/UX and accessibility	82	A-	Clear structure, strong CTAs, labelled form, skip link, reduced-motion CSS; motion not fully honored by JavaScript.
Responsiveness	80	B+	Correct viewport and engineered breakpoints; on-device verification still needed.
Content and credibility	68	B-	Excellent voice and service depth; anonymized case studies, no team, logos, testimonials, or certifications.
Visual richness and brand	90	A	Distinctive, polished, and cohesive; watch animation density and its performance cost.
Overall	78	B+	Strong foundation and design; credibility and performance are the levers to convert it into growth.

Scores are a considered judgment, not a single automated number. Where a result could not be measured directly, that is stated in the method note and reflected as a cautious rather than a low score.

Method and scope

The site was crawled live in a real Chromium browser on 11 July 2026. Pages reviewed include the English home page, the service pages (web applications and its siblings), the work index and case-study detail pages such as the operations platform, careers, how we build, the contact section and footer, and the Vietnamese home page.

Signals were gathered from the rendered DOM, from meta tags and JSON-LD extracted with JavaScript, from robots.txt, sitemap.xml, and an llms.txt probe, from the in-browser Performance API (navigation timing, paint, resource weight, and long tasks), and from accessibility checks covering form labels, reduced-motion CSS, text contrast sampling, and heading markup. Responsive behavior was read from the site's own media-query breakpoints and viewport configuration.

Limitations, stated plainly

- Google PageSpeed Insights returned a quota error during the audit, so lab Lighthouse category scores could not be pulled. Performance findings rely on the in-browser Performance API and observed behavior.
- The site has no Chrome UX Report field data yet, which is normal for a low-traffic site, so real-user Core Web Vitals are not yet observable. Targets are stated instead of measured.
- The browser tool clamped the render viewport near 1406 px, so phone and tablet layouts were assessed from the site's breakpoints and structure rather than pixel-level device rendering. An on-device pass is recommended and included as a Phase 2 item.

Findings by dimension

Technical SEO

This is the strongest part of the site. The English home page carries a unique title of 37 characters (Turn Your Will Into Real - CyberSkill) and a 153-character meta description that leads with the value proposition. The canonical tag is self-referencing, the robots meta is index and follow, the charset is UTF-8, and the viewport is set to width equals device-width. Open Graph and a summary_large_image Twitter card are present, with a dynamically generated Open Graph image.

Multilingual signals are handled correctly. The page declares hreflang for en, vi, and x-default, and those alternates are reciprocal. robots.txt is present, allows crawling, and points to the sitemap. sitemap.xml lists 20 clean URLs that cover every page in both English and Vietnamese, including the service and case-study pages.

The structured data is the standout. A JSON-LD Organization graph declares the legal name CyberSkill Software Solutions Consultancy and Development JSC, a founding date of 2020, the slogan, the email and phone, a full postal address, and a DUNS identifier. Separate FAQPage and OfferCatalog blocks describe the frequently asked questions and the services. This entity foundation is ahead of most boutique competitors.

Gaps

- The Vietnamese page keeps the English title tag (Turn Your Will Into Real - CyberSkill) even though its meta description and H1 are translated. Localize the title and meta title for the vi routes.
- robots.txt includes a non-standard Host directive and takes no explicit position on AI crawlers. Allowing all already permits them, but a deliberate stance is worth setting (see GEO).

On-page SEO

Each page uses a single H1 and a logical heading order, and section headings are descriptive. The home page holds roughly 944 words, and the service and case-study pages carry focused, benefit-led copy with clear sub-headings such as who it is for, what you get, and how we build it, or the challenge, what we did, and the outcome.

There is one specific and worthwhile fix. The large animated section headings on the home page (for example the arc of a wish, why teams work with us, and what we build) are built with a per-word kinetic-text component. Each word sits in its own span marked aria-hidden, and the heading carries an aria-label with the correct phrase. Screen readers therefore read the heading correctly. The side effect is that the DOM text serializes without spaces, so a search engine or text extractor reading those headings sees run-together words such as Thearcofawish. Add spacing so the machine-readable text matches the visible text, while keeping the aria-label.

- Only two image elements exist on the home page and both lack alt text. Most visuals are CSS or WebGL, so image-alt exposure is small, but the logo and any content images should carry meaningful alt text.
- Some routes return thin extracted body text because content is component-driven. The meaningful copy is largely present in the server-rendered HTML, which is good, but more indexable long-form content would help rankings.

GEO, or generative engine optimization

GEO is how well the site is set up to be surfaced and cited by AI answer engines such as ChatGPT, Perplexity, Gemini, and Claude. CyberSkill starts from a good place. The Organization, Service, and FAQPage schema, plus a real FAQ that answers what CyberSkill does, where it is based, how to start a project, how fast it replies, and whether it works with international clients, is exactly the kind of clean, extractable question-and-answer content these engines lift. The entity is clear and consistent across pages.

The gaps are specific. There is no llms.txt (the probe returns a 404), there is little quotable, statistic-rich or citation-rich content that AI engines tend to reward, there is no insights or blog content to be cited, and robots.txt takes no explicit position on the AI crawlers GPTBot, ClaudeBot, PerplexityBot, and Google-Extended.

It is worth being honest about llms.txt. It is an emerging proposal, adoption sits around ten percent of domains, and the major engines mostly ignore it today; Google has said it will not support it. It is low-cost insurance and useful for developer-tool agents, but it is not a proven ranking lever. The reliable GEO levers here are more quotable and cited content, expanded schema, clear entity definitions, and staying crawlable to the AI bots.

Performance and Core Web Vitals

Delivery is fast. The server responds in roughly 65 to 75 ms, which points to edge hosting, and the server-rendered content paints quickly, around one second on the content pages tested. On a warm cache the load event completes near 2.1 seconds, and repeat-visit transfers are close to zero thanks to strong caching.

The cost sits in the client bundle and the animation. The home page ships roughly 0.74 MB of decoded JavaScript across 19 script files, plus about 0.11 MB of CSS, and it runs a continuous WebGL genie and particle system that animate constantly. At least one long main-thread task of about 197 ms was observed. Heavy, always-on motion is the main risk to interaction responsiveness (INP) and to battery and CPU on the lower-end phones that are common in target markets.

Two unknowns should be closed. PageSpeed Insights was quota-blocked during the audit, and the site has no field data yet, so real-user LCP, INP, and CLS are not verifiable right now. The targets to hold at the 75th percentile of real users are LCP at or under 2.5 seconds, INP at or under 200 ms, and CLS at or under 0.1. Note also that the web-applications service page promises interfaces with Core Web Vitals kept in the green, so measuring and publishing the site's own vitals is both a credibility and a consistency matter.

UI/UX and accessibility

The experience is clear and confident. The top navigation covers Services, Work, Team, Careers, and Contact, with language and theme toggles. A scroll-driven narrative runs down the home page with a labelled progress rail, and the primary calls to action (Start my project, Talk to Lumi, and Make a wish) repeat at natural decision points. The FAQ uses an accordion, and the contact section offers both a conversational path through Lumi and a classic form.

Accessibility groundwork is present. The classic contact form associates proper labels with the name, email, company, and message fields and includes a consent checkbox, which is friendly to data-protection expectations. A skip-to-content link is present, and the stylesheet includes 34 prefers-reduced-motion rules.

Gaps

- Despite the reduced-motion CSS, the JavaScript and WebGL genie and particle animations kept running while the test environment requested reduced motion. Gate the canvas and particle animation on the same preference so decorative motion actually stops for users who ask for it.
- The classic contact form marks no fields as required in the markup, including name, email, and message. Add validation so incomplete submissions are caught.
- A leading input named website may be a spam honeypot. Confirm it is hidden from real users and from assistive technology.
- Body text contrast is high (cream on near-black), which is good. Verify the smaller muted labels and the gold-on-dark micro-labels against WCAG AA contrast.

Responsiveness

The layout is built to adapt rather than retrofitted. The viewport meta is correct and the stylesheet defines an engineered set of breakpoints (max-width rules at 440, 520, 760, 960, and 1023 pixels, and min-width rules at 720, 760, 860, and 900). The three-up card grids, the marquee, and the hero all carry breakpoint handling.

The verification gap is honest and worth naming. The interactive browser tool clamped the render viewport near 1406 px, so true phone and tablet rendering could not be exercised during this audit. Confirm tap-target sizes, the mobile navigation pattern, and that the heavy hero animation degrades gracefully on small screens as part of a device-lab pass.

Content and credibility

The writing is a real strength. The messaging is clear, specific, and in good English, with a coherent metaphor carried across the hero, the process, the services, and the contact section: a wish becomes working software, granted by Lumi. The service pages set out audience, deliverables, and method, and the voice is honest and outcome-oriented, with lines such as we name trade-offs out loud and senior engineers own the work end to end.

Credibility is the biggest single lever for global client and partner acquisition, and it is where the site is thinnest.

- Case studies are anonymized, with references to the team and a member app rather than named clients, and with no logos or quantified outcomes. Named case studies with measurable results are table stakes for consultancy sites and the highest-impact conversion asset.
- The Team navigation item points to a one-line company statement with no named people, photos, or credentials, which undercuts the site's own claim that senior engineers own the project.
- There are no client logos, testimonials, third-party ratings such as Clutch or G2, certifications such as ISO 27001 or SOC 2, or awards, which are the proof enterprise and international buyers look for.
- There is no blog or insights section, so there is no authority content to rank, to be cited by AI engines, or to demonstrate expertise.
- There is no stated engagement model or pricing signal. Strong boutiques publish engagement types, a typical project size, or a rate range to pre-qualify leads and reduce friction.

Visual richness and brand

The identity is distinctive and cohesive. A disciplined dark and gold palette, a custom 3D golden genie that reacts to interaction, particle and scroll effects, generous whitespace, and confident display typography combine into a modern, design-led impression that stands out from template-driven competitors. This is a genuine asset for differentiation as the company scales.

There are two watch-outs. The density of always-on motion competes with the content for attention and carries the performance cost described above, so it should be tuned rather than removed. And because a theme toggle exists, the light theme should be held to the same standard of polish as the dark default. Confirm that the mascot and imagery are original works, which matters more as the brand travels internationally.

Benchmark

The audit is judged against current best practice for software consultancy sites (2026) and against two reference sites: thoughtbot, a globally recognized product and engineering boutique, and Saigon Technology, a reputable Vietnam-based software firm that sells to overseas clients.

Against the technical baseline

On the modern technical checklist (unique titles and descriptions, canonical, XML sitemap, robots.txt, HTTPS, reciprocal hreflang, Organization, Service, and FAQ schema, and mobile-first delivery) CyberSkill meets or exceeds most items. Its structured data and bilingual setup are ahead of the typical boutique.

Against Core Web Vitals targets

The current position is unproven and at risk on INP because of the constant animation. This is the gap to measure and close first, since Core Web Vitals are judged on real-user field data at the 75th percentile.

Core Web Vital	Good threshold (p75 field data)	CyberSkill position
LCP (loading)	2.5 s or less	Likely fine on content pages; home hero and heavy JavaScript need verification.
INP (responsiveness)	200 ms or less	At risk from always-on WebGL and particle animation; not yet measured.
CLS (visual stability)	0.1 or less	No obvious shift observed; confirm with field data.

Note: INP replaced First Input Delay as a Core Web Vital in March 2024, so responsiveness now reflects every interaction in a visit, not just the first. That makes the animation load more important to manage than it would have been under the old metric.

Against trust-signal table stakes

This is where CyberSkill trails the benchmarks. The elements below are present on strong consultancy sites and expected by the buyers CyberSkill wants.

Trust element	Best-practice standard	CyberSkill today
Case studies	Named client or sector, problem, approach, quantified outcome	Anonymized, no metrics or logos
Testimonials	Full name, title, company, specific result	None found
Team	Named people, photos, roles, credentials	One-line company statement, no people
Client logos	Prominent, real, current	None
Third-party proof	Clutch, G2, or similar verified ratings	None
Certifications	ISO 27001, SOC 2 for enterprise and	None displayed

Trust element	Best-practice standard	CyberSkill today
	regulated buyers	
Engagement model	Engagement types or a rate or project-size signal	Not stated

What to emulate from the reference sites

- **thoughtbot (global boutique):** deep, honest case studies; a strong public content and playbook layer plus a long-running engineering blog that doubles as SEO and answer-engine fuel and as proof of expertise; and a transparent engagement model with published rates and minimums. This content-led credibility is replicable without an enterprise budget.
- **Saigon Technology (Vietnam):** the local-market pattern for winning overseas clients, with prominent certification and compliance badges, explicit global-client proof, and clear service framing.

Multi-phase remediation plan

The plan is ordered by impact and effort. Phase 1 is a set of quick, high-return fixes that can ship in days to two weeks. Phase 2 is the credibility and performance work that takes one to two months. Phase 3 is the longer-term investment that compounds over a quarter and beyond. Each item states what to change, why it matters, a rough effort, and a priority.

Phase 1: quick wins (days to about two weeks)

Action	Why it matters (impact)	Effort	Priority
Fix animated heading spacing Add word spacing to the kinetic-text H2s so DOM text matches the visible text; keep the aria-label.	Restores how search engines and AI answer engines parse the main homepage headings.	Low	High
Localize Vietnamese titles Translate the title and meta title on the vi routes, not only the descriptions.	Correct Vietnamese search results and accurate listings for the local market.	Low	High
Honor reduced motion in JavaScript Gate the WebGL genie and particle system on prefers-reduced-motion to match the CSS.	Accessibility fix and a quick responsiveness and CPU win for sensitive users.	Low to medium	High
Measure real performance Run PageSpeed Insights and Lighthouse with your own key and add a field-data tool (web-vitals or CrUX).	Replaces the current unknown with numbers before any optimization work.	Low	High
Surface trust you already hold Add two or three named client logos or attributed testimonials, plus the registration you hold, near the top of the home and contact areas.	Immediate credibility for prospects and partners.	Low to medium	High
Add image alt text Give the logo and any content images meaningful alt text.	Accessibility and image search coverage.	Low	Medium
Publish llms.txt and set an AI-crawler stance Add a root llms.txt and state a deliberate position on GPTBot, ClaudeBot, PerplexityBot, and Google-Extended in robots.txt.	Low-cost answer-engine insurance and developer-tool discoverability; keeps you retrievable.	Low	Medium
Validate the contact form Mark name, email, and message as required, add validation, and	Fewer incomplete or spam leads and cleaner accessibility.	Low	Medium

Action	Why it matters (impact)	Effort	Priority
confirm the website honeypot is hidden.			

Phase 2: medium-term improvements (about one to two months)

Action	Why it matters (impact)	Effort	Priority
Rebuild case studies with proof Name the client or sector, state the problem, approach, and quantified results (hours saved, error rate, uptime, revenue, load time), and add screenshots.	The single biggest conversion lever for a consultancy site.	Medium	High
Build a real team page Named senior engineers with roles, short bios, credentials, and links.	Matches the senior-engineers-own-the-work promise and builds buyer trust.	Medium	High
Cut and split the JavaScript Code-split the animation and 3D so they load after content and only where used, defer non-critical scripts, set a performance budget, and pause off-screen animation.	Protects Core Web Vitals and the low-end mobile experience.	Medium to high	High
Add third-party proof Create a Clutch and/or G2 profile, gather verified reviews, and display badges.	Independent social proof that international buyers trust more than on-site quotes.	Medium	Medium
Expand structured data Add BreadcrumbList on subpages, Service schema per service page, Article schema when the blog launches, and Review or AggregateRating once testimonials are verifiable.	Richer search results and better AI attribution.	Medium	Medium
Start an insights section Publish a few strong, quotable articles with statistics and citations tied to your services and market.	Authority content for search, answer-engine citation, and expertise signals.	Medium	Medium
State an engagement model Show engagement types and a typical project size or rate range.	Pre-qualifies leads and reduces friction before the first call.	Low to medium	Medium
Run an on-device audit Test real phones and tablets, keyboard-only navigation, a screen reader, and WCAG AA contrast on the muted labels.	Closes the responsiveness and accessibility verification gap.	Medium	Medium

Phase 3: longer-term strategic investments (a quarter and beyond)

Action	Why it matters (impact)	Effort	Priority
Build a content and authority engine A durable library of guides and process artifacts and a regular publishing cadence, in the spirit of a public playbook.	Compounding search and answer-engine visibility and a steady lead source; the long-term investment with the most compounding return.	High	Medium
Pursue certifications Work toward ISO 27001 and/or SOC 2 Type II, then display them.	Unlocks larger, enterprise, and regulated deals, especially international ones.	High	Medium
Monitor performance continuously Stand up field-data dashboards and alerts that catch Core Web Vitals regressions.	Keeps the in-the-green promise true as the site grows.	Medium	Medium
Deepen a GEO content program Publish quotable, cited, statistic-rich answers to the questions buyers ask AI engines about software consultancies in Vietnam and your specialties.	Positions CyberSkill to be the cited answer in AI results.	Medium to high	Medium
Expand international readiness Add region-specific messaging or pages and localized proof, and keep hreflang reciprocal as pages grow.	Directly supports the global-scaling goal.	High	Low to medium

Closing note

The design and engineering quality on cyberskill.world is already a differentiator, and the technical SEO and structured data are ahead of most peers. The next gains come from proof and measurement: name the clients and the numbers, show the team, verify and protect Core Web Vitals, and ship a short list of precise search and answer-engine fixes. Do that and the site will convert more of the attention its design earns into the clients and partners the business is built to grow.